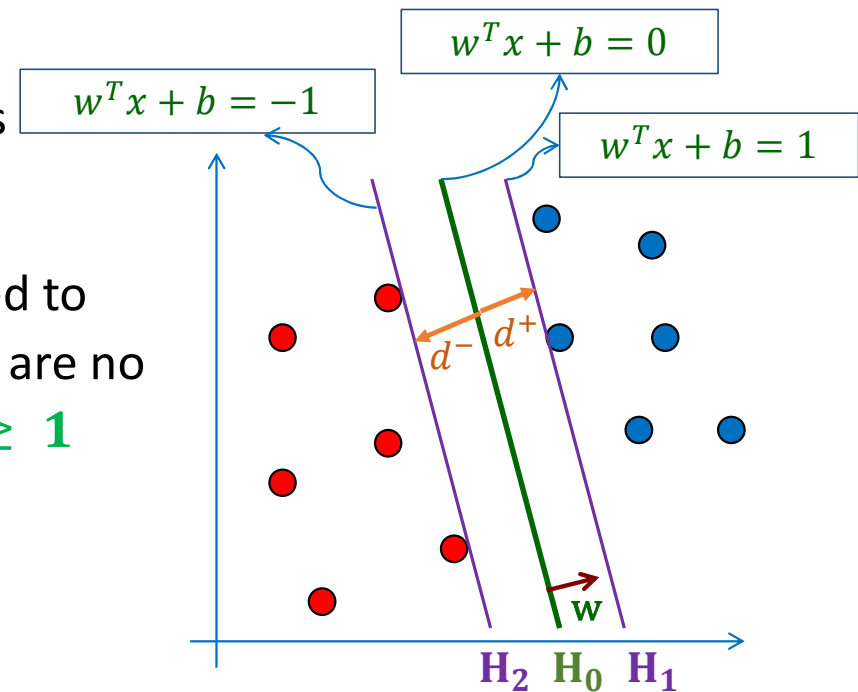


Non Linear SVM

Kernel Trick

Maximizing the Margin

- We want a classifier (linear separator) with as big a margin as possible.
- In order to maximize the margin, we thus need to minimize $\|\mathbf{w}\|$. With the condition that there are no datapoints between H_1 and H_2 : $\mathbf{y}_i (\mathbf{w}^T \mathbf{x}_i) \geq 1$
- Minimize $J(\mathbf{w}) = \frac{1}{2} \mathbf{w}^T \mathbf{w}$, subject to:
 $\forall_i, \mathbf{y}_i (\mathbf{w}^T \mathbf{x}_i + b) \geq 1$
- Constrained quadratic optimization problem solved by the **Lagrangian multiplier method**

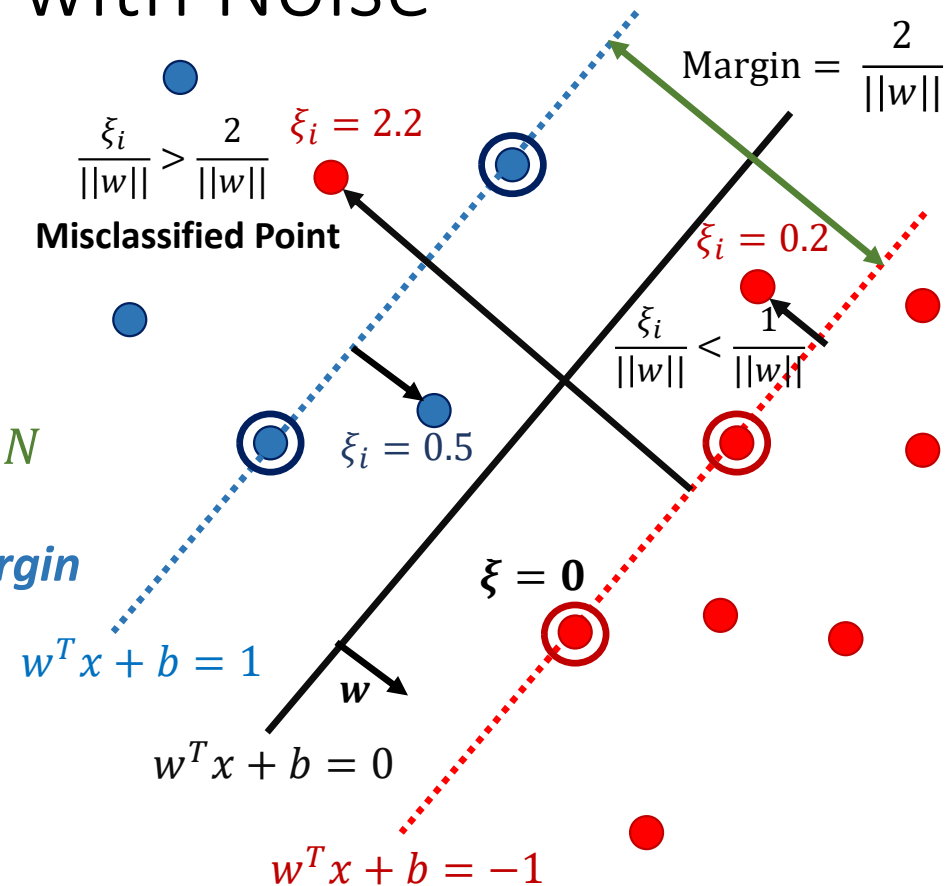


Learning Maximum Margin with Noise

- The error terms $\xi_N'^S$ are incorporated into our optimization problem by:

$$\left\{ \begin{array}{l} \min_{w,b} ||w||^2 + C \sum_{i=1}^N \xi_i \\ \text{such that } y_i(w^T x_i + b) \geq 1 - \xi_i, \quad i = 1, \dots, N \end{array} \right.$$

- The solution to this problem is called **soft margin support vector classification**



SVM Solution

Involves computing the *inner products* $\mathbf{x}_i^T \mathbf{x}_j$ between all training points

1. Maximize: $Q(\boldsymbol{\alpha}) = \sum_{i=1}^N \alpha_i - \frac{1}{2} \sum_{i=1}^N \sum_{j=1}^N \alpha_i \alpha_j y_i y_j \mathbf{x}_i^T \mathbf{x}_j$, subject to $0 \leq \alpha_i \leq C \quad \forall i$,
 $\sum_{i=1}^N \alpha_i y_i = 0$
- $\mathbf{w} = \sum_{i=1}^N \alpha_i y_i \mathbf{x}_i$; (very few α_i s are non-zero: support vectors. The sum is therefore only to be over the support vectors)
- $\mathbf{b} = \mathbf{y}_K(1 - \xi_K) - \mathbf{w}_K^T \mathbf{x}_K = \mathbf{y}_K(1 - \xi_K) - \sum_{i=1}^N \alpha_i y_i \mathbf{x}_i^T \mathbf{x}_K$ with $K = \arg \max_i \alpha_i$

Note: Classification:

$$\mathbf{f}(\mathbf{x}_t) = \mathbf{w}^T \mathbf{x}_t + \mathbf{b} = \sum_{i=1}^N \alpha_i y_i \mathbf{x}_i^T \mathbf{x}_t + b$$

Relies on an *inner product* between the test point \mathbf{x}_t and the support vectors \mathbf{x}_i

Theoretical Justification for Maximum Margin

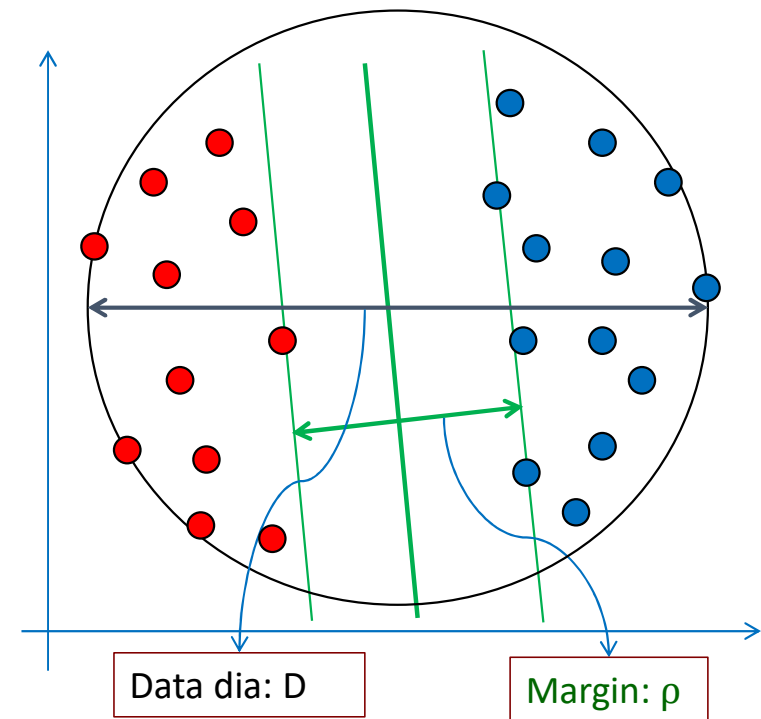
- V.N. Vapnik and A.Ya. Červonenkis quantified complexity
 - Higher the VC-dimension (h), more complex the classifier
 - Bound on Expected Loss:

$$R_{tst}(\alpha) = R_{trn}(\alpha) + f(h, N)$$

$$h \leq \min \left\{ d, \left\lceil \frac{1}{m^2} \right\rceil \right\} + 1,$$

where d is the dimensionality. $m = \rho/D$ is the relative margin, with ρ as the margin, and D as the diameter of the smallest sphere that encloses all of the training examples.

- Implication: If ρ/D is high, VC dimension is low and expected error is low, regardless of the dimensionality d
- Complexity of the classifier is kept small regardless of dimensionality.

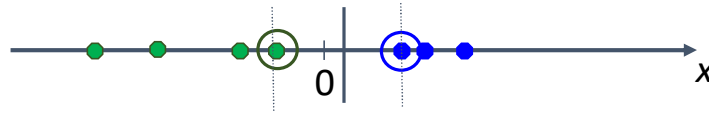


Linear SVMs: Overview

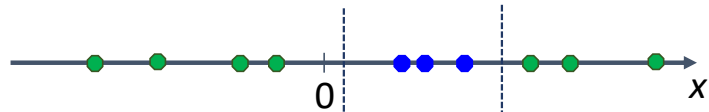
- Convex Optimization guarantees the global optimum.
- Support vectors are automatically identified.
- Use of max margin optimizes test accuracy
 - One of the best classifiers, given a feature representation
 - SVMs works well even with fewer training samples.
 - Note: Minimizes overfitting
 - Does not scale to huge datasets
 - GD based approaches work better in such cases.

Non-linear SVMs

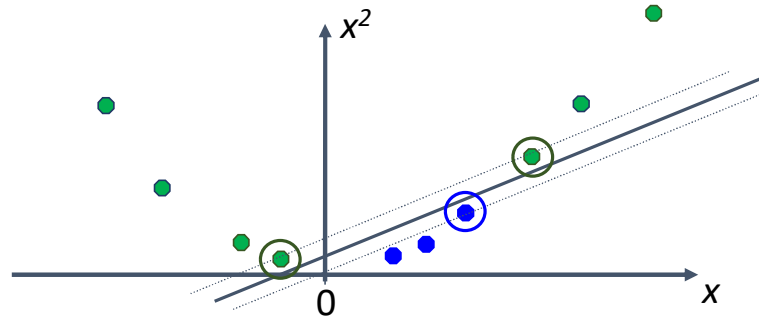
- Datasets that are linearly separable with some noise work out great with Linear SVMs:



- But what should be done if the dataset is just not linearly separable.

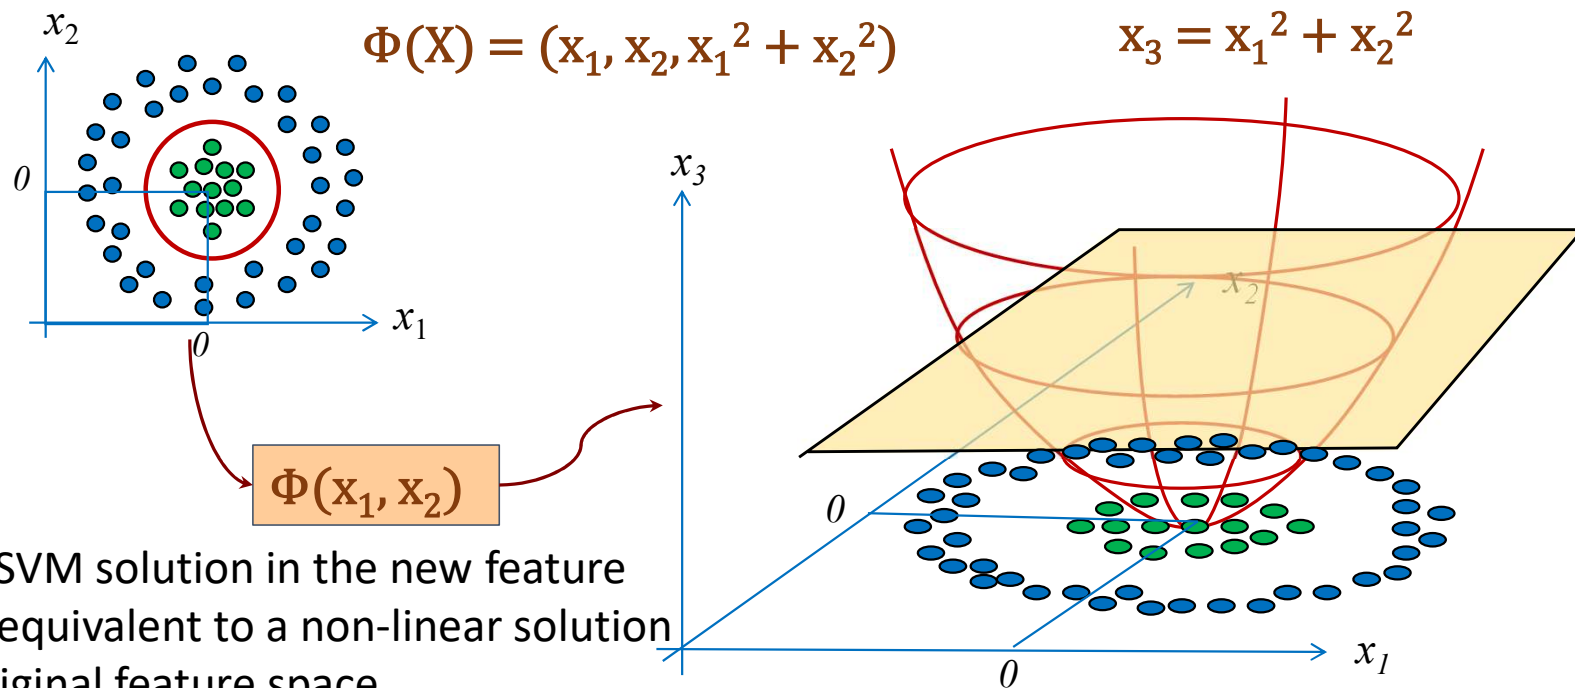


- Apply a non-linear transformation, to the feature space such that the samples become linearly separable



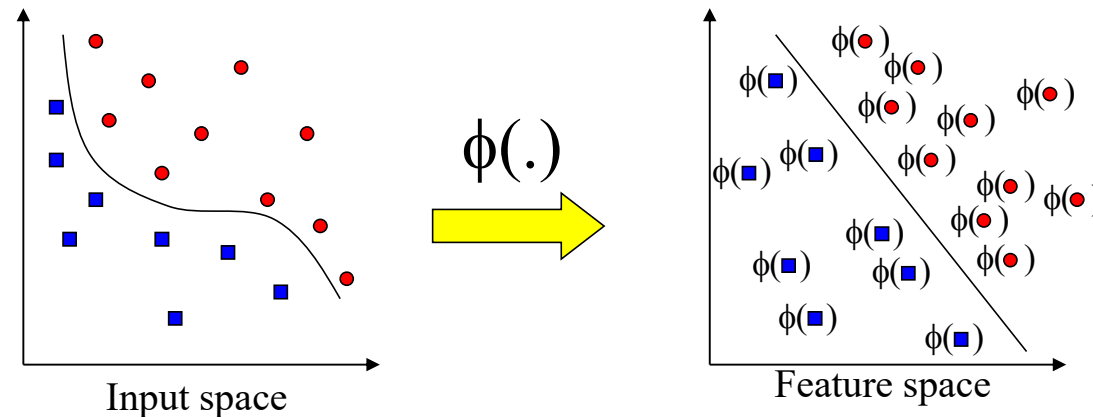
$$\Phi_k = (x_k, x_k^2)$$

Circular Boundary



- A linear SVM solution in the new feature space is equivalent to a non-linear solution in the original feature space.
- The mapping, $\Phi(X)$, is however unknown. Depends on the distribution of points in the feature space. It can be a complex mapping in general

Transforming the Data



Note: feature space is of higher dimension than the input space in practice

- Computation in the feature space can be costly because it is high dimensional
 - The feature space is typically infinite-dimensional!
- Question: Can we find the SVM solution without knowing $\Phi(\mathbf{X})$?
- The kernel trick comes to rescue

Quadratic Basis Function

- Let there be a mapping from $\mathbf{x} \rightarrow \Phi(\mathbf{x})$

$$\mathbf{x} = \begin{bmatrix} x_1 \\ x_2 \\ x_3 \end{bmatrix} \quad \text{Choose } \Phi(\mathbf{x}) \text{ to be pairwise monomial terms,} \\ \Phi(\mathbf{x}) \in \mathbb{R}^{m^2}, \text{ with } \mathbf{x} \in \mathbb{R}^m$$


$$\Phi(\mathbf{x}) \rightarrow \begin{bmatrix} x_1 x_1 \\ x_1 x_2 \\ x_1 x_3 \\ x_2 x_1 \\ x_2 x_2 \\ x_2 x_3 \\ x_3 x_1 \\ x_3 x_2 \\ x_3 x_3 \end{bmatrix} \quad \Phi(\mathbf{z}) \rightarrow \begin{bmatrix} z_1 z_1 \\ z_1 z_2 \\ z_1 z_3 \\ z_2 z_1 \\ z_2 z_2 \\ z_2 z_3 \\ z_3 z_1 \\ z_3 z_2 \\ z_3 z_3 \end{bmatrix}$$

$$\mathbf{x} \rightarrow \Phi(\mathbf{x}) \rightarrow \begin{bmatrix} x \\ x^2 \\ x^3 \end{bmatrix}$$

$$\begin{bmatrix} x_1 \\ x_2 \end{bmatrix} \rightarrow \Phi(\mathbf{x}) \rightarrow \begin{bmatrix} x_1 \\ x_2 \\ x_1^2 \\ x_2^2 \\ x_1 x_2 \\ \dots \end{bmatrix}$$

$$\Phi^T(x) \Phi(z) = [x_1 x_1 \ x_1 x_2 \ x_1 x_3 \ \dots \ x_3 x_3] \begin{bmatrix} z_1 z_1 \\ z_1 z_2 \\ z_1 z_3 \\ z_2 z_1 \\ z_2 z_2 \\ z_2 z_3 \\ z_3 z_1 \\ z_3 z_2 \\ z_3 z_3 \end{bmatrix} = \sum_{i=1}^3 \sum_{j=1}^3 (x_i x_j) (z_i z_j)$$

SVM Solution


1. Maximize: $Q(\alpha) = \sum_{i=1}^N \alpha_i - \frac{1}{2} \sum_{i=1}^N \sum_{j=1}^N \alpha_i \alpha_j y_i y_j \mathbf{x}_i^T \mathbf{x}_j$, 

subject to $0 \leq \alpha_i \leq C \quad \forall i, \quad \sum_{i=1}^N \alpha_i y_i = 0$

- $\mathbf{w} = \sum_{i=1}^N \alpha_i y_i \mathbf{x}_i = \sum_{i=1}^N \alpha_i y_i \Phi(\mathbf{x}_i)$

Note: Classification:

$$\mathbf{f}(\mathbf{x}_t) = \mathbf{w}^T \mathbf{x}_t + b = \sum_{i=1}^N \alpha_i y_i \mathbf{x}_i^T \mathbf{x}_t + b$$



- We need to do $\frac{N(N+1)}{2} \approx \frac{N^2}{2}$ dot products to calculate $\Phi^T(\mathbf{x}_i)\Phi(\mathbf{x}_j)$
- Each dot product further requires m^2 calculations
- The whole calculation will cost $\frac{N^2 m^2}{2}$

Quadratic Basis Function

$$\mathbf{x} = \begin{bmatrix} x_1 \\ x_2 \\ x_3 \end{bmatrix} \quad \Phi(\mathbf{x}) \rightarrow \begin{bmatrix} x_1 x_1 \\ x_1 x_2 \\ x_1 x_3 \\ x_2 x_1 \\ x_2 x_2 \\ x_2 x_3 \\ x_3 x_1 \\ x_3 x_2 \\ x_3 x_3 \end{bmatrix} \quad \Phi(\mathbf{z}) \rightarrow \begin{bmatrix} z_1 z_1 \\ z_1 z_2 \\ z_1 z_3 \\ z_2 z_1 \\ z_2 z_2 \\ z_2 z_3 \\ z_3 z_1 \\ z_3 z_2 \\ z_3 z_3 \end{bmatrix}$$

Choose $\Phi(\mathbf{x})$ to be pairwise monomial terms,
 $\Phi(\mathbf{x}) \in \mathbb{R}^{m^2}$, with $\mathbf{x} \in \mathbb{R}^m$

$$\begin{aligned} \Phi^T(\mathbf{x})\Phi(\mathbf{z}) &= [x_1 x_1 \ x_1 x_2 \ x_1 x_3 \ \dots \ x_3 x_3] \begin{bmatrix} z_1 z_1 \\ z_1 z_2 \\ z_1 z_3 \\ z_2 z_1 \\ z_2 z_2 \\ z_2 z_3 \\ z_3 z_1 \\ z_3 z_2 \\ z_3 z_3 \end{bmatrix} \\ &= \sum_{i=1}^3 \sum_{j=1}^3 (x_i x_j)(z_i z_j) \end{aligned}$$

$O(m^2)$

- Just out of casual, innocent, interest, let's look at another function of \mathbf{x} and \mathbf{z} :

$$(\mathbf{x}^T \mathbf{z})^2 = \left(\sum_{i=1}^3 x_i z_i \right) \left(\sum_{j=1}^3 x_j z_j \right) = \sum_{i=1}^3 \sum_{j=1}^3 x_i z_i x_j z_j = \sum_{i=1}^3 \sum_{j=1}^3 (x_i x_j)(z_i z_j)$$

Both are same, but
 $(\mathbf{x}^T \mathbf{z})^2$ is only $O(m)$ to
 compute

Quadratic Form

$$\mathbf{x} \rightarrow \Phi(\mathbf{x})$$

$$m = 3$$

$$\begin{bmatrix} x_1 \\ x_2 \\ x_3 \end{bmatrix} \rightarrow \Phi(\mathbf{x}) \rightarrow \begin{bmatrix} 1 \\ \sqrt{2}x_1 \\ \sqrt{2}x_2 \\ \sqrt{2}x_3 \\ x_1^2 \\ x_2^2 \\ x_3^2 \\ \sqrt{2}x_1x_2 \\ \sqrt{2}x_1x_3 \\ \sqrt{2}x_2x_3 \end{bmatrix} = O(10) = O(2m + \frac{m(m-1)}{2} + 1) \approx O\left(\frac{m^2}{2}\right)$$

$$\Phi(\mathbf{x}) \cdot \Phi(\mathbf{z}) = 1 + 2 \sum_{i=1}^m x_i z_i + \sum_{i=1}^m x_i^2 z_i^2 + \sum_{i=1}^m \sum_{j=i+1}^m 2x_i x_j z_i z_j$$

$$(x \cdot z + 1)^2 = (x \cdot z)^2 + 2x \cdot z + 1 = \left(\sum_{i=1}^m x_i z_i \right)^2 + 2 \sum_{i=1}^m x_i z_i + 1 = \sum_{i=1}^m \sum_{j=1}^m x_i z_i x_j z_j + 2 \sum_{i=1}^m x_i z_i + 1$$

$$= 1 + 2 \sum_{i=1}^m x_i z_i + \sum_{i=1}^m x_i^2 z_i^2 + \sum_{i=1}^m \sum_{j=i+1}^m 2x_i x_j z_i z_j \quad O(m)$$

For a Cubic Kernel

- If the original Space is 3-dimensional:

$$\mathbf{K}(\mathbf{X}, \mathbf{Z}) = (\mathbf{X} \cdot \mathbf{Z})^3 = (x_1 z_1 + x_2 z_2 + x_3 z_3)^3$$

- Equivalent to working in a 10-dimensional space
- Kernel: 5(3+2):× and 2:+
- $\Phi(X) = 13: \times$, $\Phi(Z) = 13: \times$,
- $\Phi(X)\Phi(Z) = 10: \times, 9: +$
- Total 36(13+13+10):× and 9:+

$$\Phi(\mathbf{X}) = \Phi\left(\begin{bmatrix} x_1 \\ x_2 \\ x_3 \end{bmatrix}\right) = \begin{bmatrix} x_1^3 \\ x_2^3 \\ x_3^3 \\ x_1^2 x_2 \\ x_1 x_2^2 \\ x_1^2 x_3 \\ x_1 x_3^2 \\ x_2^2 x_3 \\ x_2 x_3^2 \\ x_1 x_2 x_3 \end{bmatrix}$$

Higher Order Polynomials

Kernel Function

Polynomial	$\Phi(\mathbf{x})$	Cost to build Q_{kl} matrix traditionally	Cost if 100 features	$\Phi(\mathbf{x}) \cdot \Phi(\mathbf{z})$ $K(\mathbf{x}, \mathbf{z})$	Cost to build Q_{kl} matrix sneakily	Cost if 100 features
Quadratic	All $m^2/2$ terms up to degree 2	$m^2 N^2 / 4$	$2,500 N^2$	$(\mathbf{x} \cdot \mathbf{z} + 1)^2$	$m N^2 / 2$	$50 N^2$
Cubic	All $m^3/6$ terms up to degree 3	$m^3 N^2 / 12$	$83,000 N^2$	$(\mathbf{x} \cdot \mathbf{z} + 1)^3$	$m N^2 / 2$	$50 N^2$
Quartic	All $m^4/24$ terms up to degree 4	$m^4 N^2 / 48$	$1,960,000 N^2$	$(\mathbf{x} \cdot \mathbf{z} + 1)^4$	$m N^2 / 2$	$50 N^2$

SVM Solution: Kernel function

1. Maximize:

$$Q(\alpha) = \sum_{i=1}^N \alpha_i - \frac{1}{2} \sum_{i=1}^N \sum_{j=1}^N \alpha_i \alpha_j y_i y_j \mathbf{K}(\mathbf{x}_i, \mathbf{x}_j)$$

2. $\mathbf{w} = \sum_{i=1}^N \alpha_i y_i \mathbf{x}_i;$

3. $\mathbf{b} = \mathbf{1} - \mathbf{w}^T \mathbf{x}_{s+} = \mathbf{1} - \sum_{i=1}^N \alpha_i y_i \mathbf{K}(\mathbf{x}_i, \mathbf{x}_{s+})$

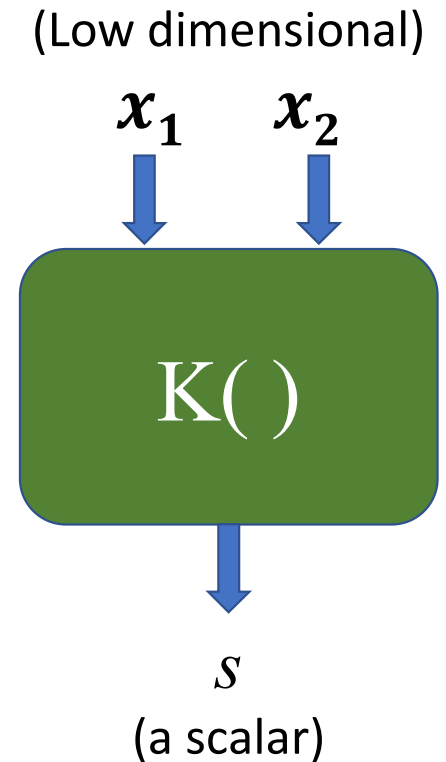
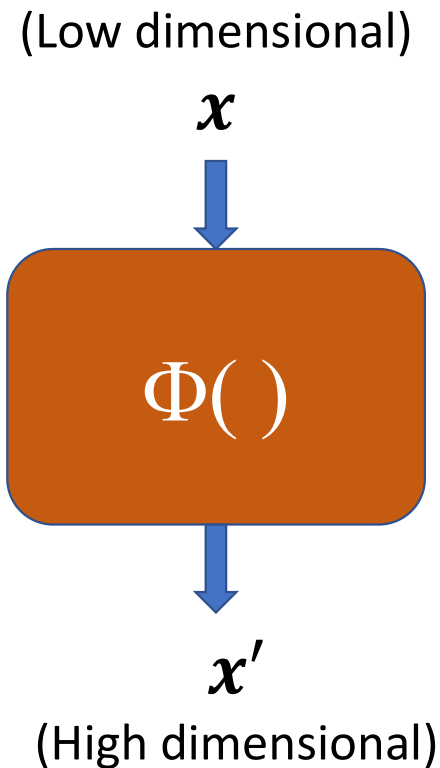
Do we know the $\mathbf{K}(\mathbf{x}_i, \mathbf{x}_j)$.?
Let us compare Φ and $\mathbf{K}()$.

Note: Classification:

$$\mathbf{g}(\mathbf{x}_t) = \mathbf{w}^T \mathbf{x}_t + \mathbf{b} = \sum_{i=1}^N \alpha_i y_i \mathbf{K}(\mathbf{x}_i, \mathbf{x}_t) + b$$

Comparing compare $\Phi()$ and $K()$

- $\Phi(\mathbf{x})$ is a complex non-linear mapping of \mathbf{x} into a high-dimensional space.
- $K(\mathbf{x}_i, \mathbf{x}_j)$ is a simple function that measures the similarity between two vectors.
- If we know the K function, we do not need Φ .
- A kernel function *implicitly* maps data to a high-dimensional space (without the need to compute each $\Phi(\mathbf{x})$ explicitly).



What Functions are Kernels?

- For some functions $K(\mathbf{x}_i, \mathbf{x}_j)$ checking that $K(\mathbf{x}_i, \mathbf{x}_j) = \Phi^T(\mathbf{x}_i)\Phi(\mathbf{x}_j)$ can be cumbersome.
- Mercer's theorem determines which functions can be used as a kernel function:

Every semi-positive definite symmetric function is a kernel

- Semi-positive definite symmetric functions correspond to a semi-positive definite symmetric Gram matrix:

$\mathbf{K} =$

$K(\mathbf{x}_1, \mathbf{x}_1)$	$K(\mathbf{x}_1, \mathbf{x}_2)$	$K(\mathbf{x}_1, \mathbf{x}_3)$...	$K(\mathbf{x}_1, \mathbf{x}_n)$
$K(\mathbf{x}_2, \mathbf{x}_1)$	$K(\mathbf{x}_2, \mathbf{x}_2)$	$K(\mathbf{x}_2, \mathbf{x}_3)$		$K(\mathbf{x}_2, \mathbf{x}_n)$
...
$K(\mathbf{x}_n, \mathbf{x}_1)$	$K(\mathbf{x}_n, \mathbf{x}_2)$	$K(\mathbf{x}_n, \mathbf{x}_3)$...	$K(\mathbf{x}_n, \mathbf{x}_n)$

✓ $K(\mathbf{x}_i, \mathbf{x}_j)$ measures the similarity or proximity between two data points \mathbf{x}_i and \mathbf{x}_j in the input space

Examples of Kernel Functions

- Linear: $K(\mathbf{x}_i, \mathbf{x}_j) = \mathbf{x}_i^T \mathbf{x}_j$

- Mapping $\Phi: \mathbf{x} \rightarrow \Phi(\mathbf{x})$, where $\Phi(\mathbf{x})$ is \mathbf{x} itself

- Polynomial of power p : $K(\mathbf{x}_i, \mathbf{x}_j) = (1 + \mathbf{x}_i^T \mathbf{x}_j)^p$

- Mapping $\Phi: \mathbf{x} \rightarrow \Phi(\mathbf{x})$, where $\Phi(\mathbf{x})$ has $d+p$ dimensions

- Gaussian (radial-basis function): $K(\mathbf{x}_i, \mathbf{x}_j) = e^{-\frac{(\|\mathbf{x}_i - \mathbf{x}_j\|)^2}{2\sigma^2}}$

- Mapping $\Phi: \mathbf{x} \rightarrow \Phi(\mathbf{x})$, where $\Phi(\mathbf{x})$ is *infinite-dimensional*: every point is mapped to a function (a Gaussian); combination of functions for support vectors is the separator.

Example

- Suppose we have 5 one-dimensional data points
 - $x_1 = 1, x_2 = 2, x_3 = 4, x_4 = 5, x_5 = 6$, with 1, 2, 6 as class 1 and 4, 5 as class 2 $\Rightarrow y_1 = 1, y_2 = 1, y_3 = -1, y_4 = -1, y_5 = 1$
- We use the polynomial kernel of degree 2
 - $K(x, y) = (xy + 1)^2$
 - C is set to 100
- We first find α_i ($i = 1, \dots, 5$) by

$$\begin{aligned} \max. \quad & \sum_{i=1}^5 \alpha_i - \frac{1}{2} \sum_{i=1}^5 \sum_{j=1}^5 \alpha_i \alpha_j y_i y_j (x_i x_j + 1)^2 \\ \text{subject to } & 0 \leq \alpha_i \leq 100, \sum_{i=1}^5 \alpha_i y_i = 0 \end{aligned}$$

Example

- By using a QP solver, we get

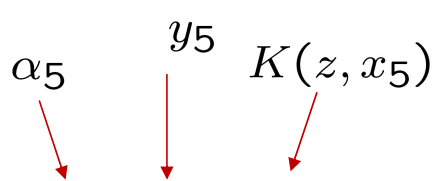
- $\alpha_1 = 0, \alpha_2 = 2.5, \alpha_3 = 0, \alpha_4 = 7.333, \alpha_5 = 4.833$

- Note that the constraints are indeed satisfied

- The support vectors are $\{x_2 = 2, x_4 = 5, x_5 = 6\}$

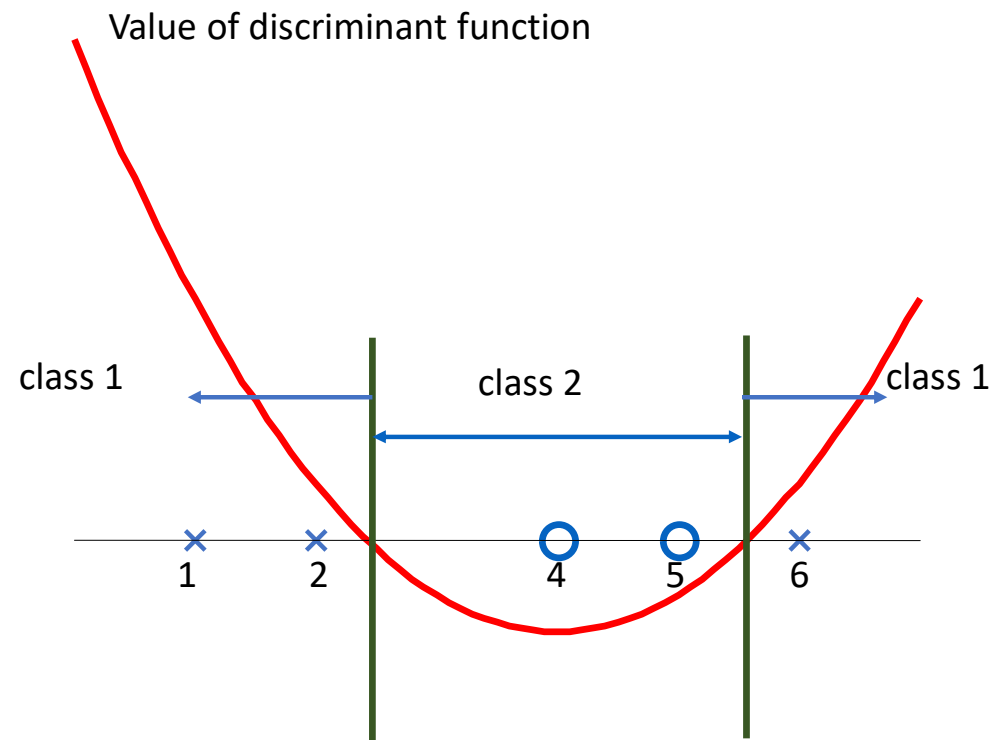
- The discriminant function is

$$f(z) = \sum_{i=1}^N \alpha_i y_i \mathbf{x}_i^T \mathbf{z} + b = 2.5 (1)(2z + 1)^2 + 7.333 (-1)(5z + 1)^2 + 4.833 (1)(6z + 1)^2 + b$$

$$= 0.6667 z^2 - 5.333 z + b$$


- b is recovered by solving $f(2)=1$ or by $f(5)=-1$ or by $f(6)=1$, as x_2 and x_5 lie on the line $\mathbf{w}^T \mathbf{x} + b = 1$ and x_4 lies on the line $\mathbf{w}^T \mathbf{x} + b = -1$
- All three give $b=9 \Rightarrow f(z) = 0.6667 z^2 - 5.333 z + 9$

Example

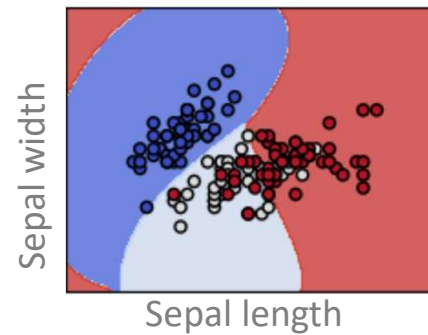


Scikit Learn Implementation

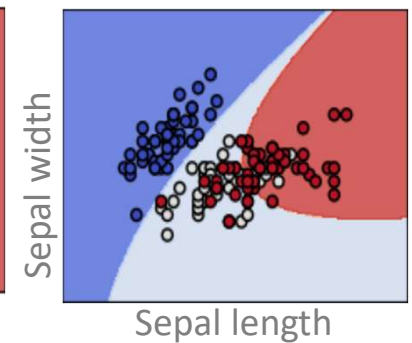
- `libsvm` based: `sklearn.svm.SVC`
- Several Kernels: linear, polynomial, rbf, sigmoid, custom.
- Other HyperParameters: C, kernel params
- Usage:

```
>>> from sklearn import svm
>>> X = [[0, 0], [1, 1]]
>>> y = [0, 1]
>>> clf = svm.SVC(kernel='rbf')
>>> clf.fit(X, y)
>>> clf.predict([[2., 2.]])
```

SVC: RBF Kernel



SVC: Poly (3) Kernel



```
>>> # get support vectors
>>> clf.support_vectors_
>>> # get support vector indices
>>> clf.support_
```

<https://scikit-learn.org/stable/modules/svm.html>

Summary

- Linear SVMs generalize well, but cannot separate non-linear data
- If features can be transformed appropriately, SVMs can learn non-linear boundaries.
- How do we find the feature transformation?
 - Use some popular Kernel functions.
- Kernels (nonlinear) SVMs are also good at generalization and can deal with non-linear data.