

Loss Function For Regression


Recap: Linear Regression

- For the n samples (x^j, y^j) , to be fitted to a line

$$y^j = \alpha_0 + \alpha_1 x_1^j + \alpha_2 x_2^j + \cdots + \alpha_r x_r^j + \varepsilon^j,$$

where ε^j is the residual for the j^{th} sample and r is the number of predictors.

- We measure goodness of fit by a loss function, defined on the residuals: $\varepsilon^j = y^j - f(x^j) = y^j - \hat{y}^j$
- A popular fit uses the least square approach that minimizes the sum of squared residuals
 - e.g., $L_{MSE} = \sum_j \varepsilon_j^2$



Loss Functions

- A loss function, L , is defined as a mapping of $f(x_i)$ with its corresponding y_i to a real number $l \in \mathbb{R}$, which captures the similarity between $f(x_i)$ and y_i

$$\varepsilon_i = y_i - f(x_i)$$

❑ Mean Bias Error: $L_{MBE} = \frac{1}{N} \sum_{i=1}^N \varepsilon_i$

❑ Mean Absolute Error: $L_{MAE} = \frac{1}{N} \sum_i |\varepsilon_i|$

- ✓ Loss is continuous and differentiable.
- ✓ Helps to identify the direction of model bias
- ✓ Errors may cancel out, leading to zero loss
- ✓ Also called Laplace or l_1 loss.
- ✓ Least affected by outliers
- ✓ Not differentiable: Can create trouble

NOTE: the loss function is to capture the difference between the actual and predicted values for a single record whereas cost functions aggregate the difference for the entire training dataset.

Loss Functions

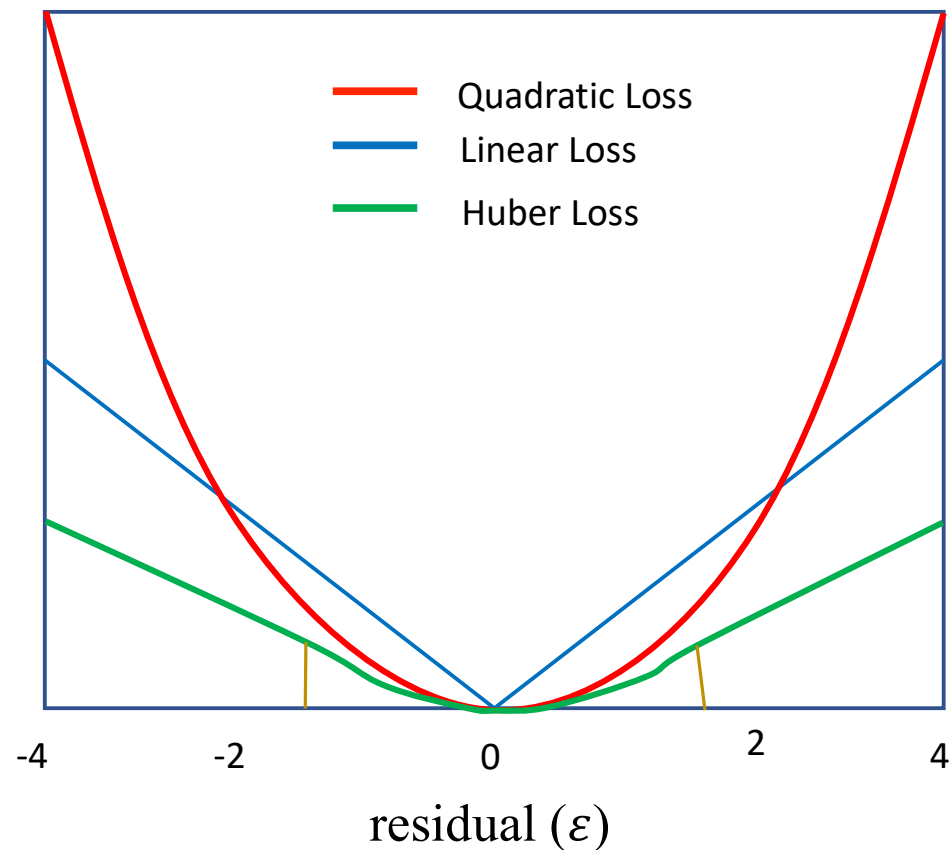
□ Mean Squared Error: $L_{MSE} = \frac{1}{N} \sum_i \varepsilon_i^2$

- ✓ Also called Quadratic or l_2 loss.
- ✓ Significantly affected by outliers
- ✓ Mathematically elegant optimization

□ Root Mean Squared: $L_{RMSE} = \sqrt{\frac{1}{N} \sum_i \varepsilon_i^2}$

- ✓ Less affected by outliers
- ✓ More expensive to compute
- ✓ Continuous and differentiable

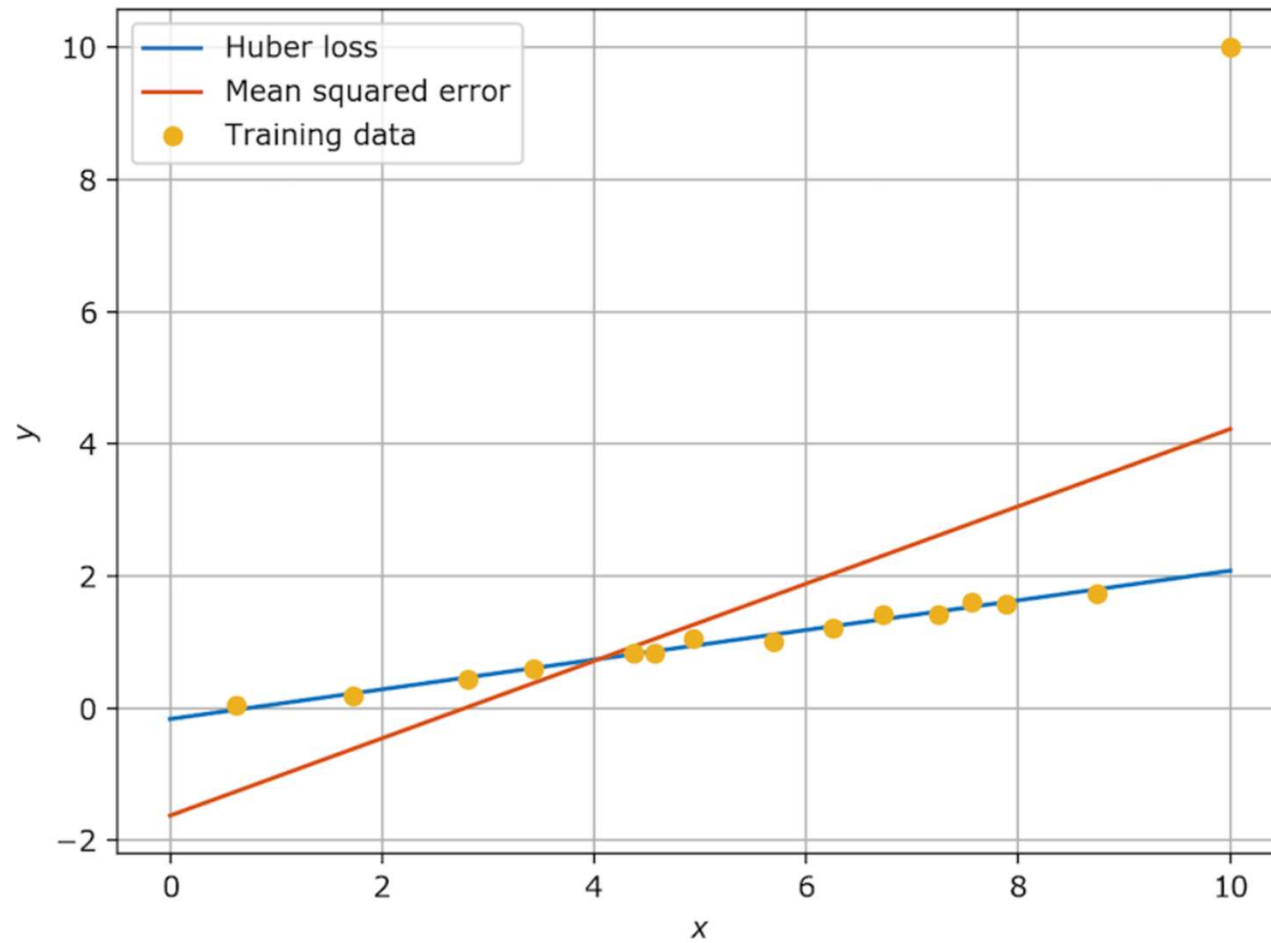
Huber loss: Robust Regression



- Combines the advantages of MSE and MAE.
- Mostly proportional to the absolute value, except for small errors, where it is proportional to the square of the error

$$L_{\delta}(\epsilon_i) = \begin{cases} \frac{1}{2} \epsilon_i^2, & |\epsilon_i| \leq \delta \\ \delta \cdot (|\epsilon_i| - \frac{1}{2} \delta), & |\epsilon_i| > \delta \end{cases}$$

- Robustness of l_1 loss
- Differentiable
- δ is adjusted during training based on what is considered an outlier




Model prediction with outlier in training data. Comparison of Huber loss and MSE.

Log-cosh loss

- Log-cosh loss is very close to the Huber loss in behavior, without the requirement of learning δ
 - linear behavior for very large values of error as **$\log(\cosh x) \rightarrow |x| - \log 2$**
 - quadratic behavior for small loss values, as **$\log(\cosh x) \rightarrow \frac{x^2}{2}$** .

$$L_{\logcosh} = \frac{1}{N} \sum_{i=1}^N \log(\cosh(f(x_i) - y_i))$$



- More computationally expensive and less customizable compared to Huber Loss

RMSLE Loss

- The Root Mean Squared Logarithmic Error (RMSLE) loss is the RMSE of the log-transformed observed value y and log-transformed predicted value \hat{y}

$$L_{RMSLE} = \sqrt{\frac{1}{N} \sum_{i=1}^N (\log(y_i + 1) - \log(f(x_i) + 1))^2}$$

- ✓ Due to the properties of the logarithm, the error between the predicted and the actual values is relative, making the RMSLE more robust to outliers

y	55	64	72	450
\hat{y}	61	59	74	102

$$RMSE = 4.655 \quad (174.047)$$

$$RMSLE = 0.076 \quad (0.741)$$

RMSLE Loss

$$L_{RMSLE} = \sqrt{\frac{1}{N} \sum_{i=1}^N (\log(y_i + 1) - \log(f(x_i) + 1))^2} = \sqrt{\frac{1}{N} \sum_{i=1}^N \left(\frac{\log(y_i + 1)}{\log(f(x_i) + 1)} \right)^2}$$

- ✓ Only considers the relative error between the Predicted and the actual value and the scale of the error is not significant
- ✓ RMSLE incurs a larger penalty for the underestimation of the Actual variable than the Overestimation.

y	100	10000
\hat{y}	80	8000

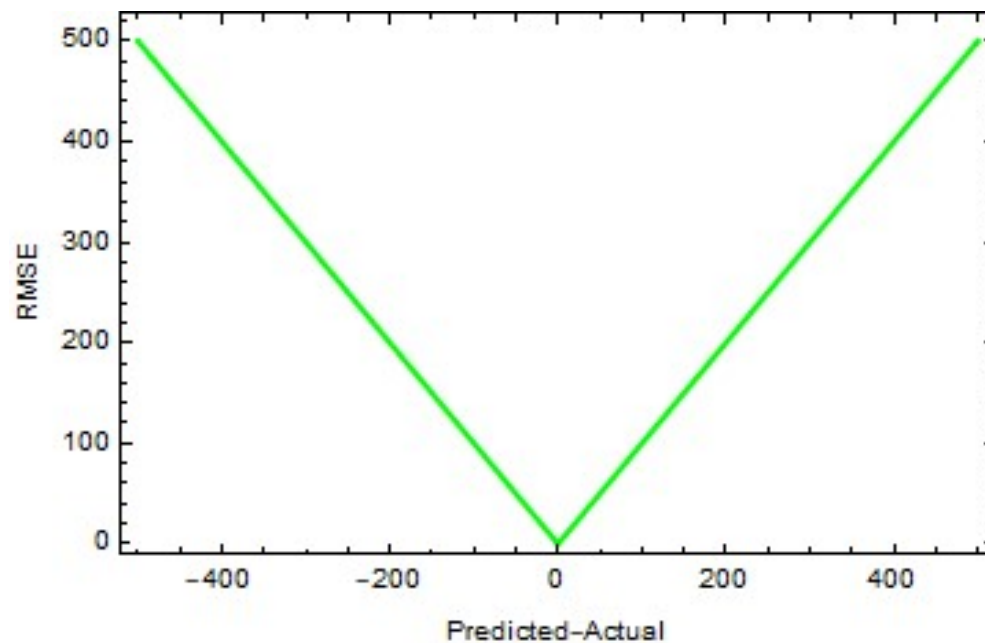

$$\begin{aligned} RMSE &= 20 & (2000) \\ RMSLE &= 0.220 & (0.223) \end{aligned}$$

y	1000	1000
\hat{y}	600	1400

$$\begin{aligned} RMSE &= 400 & (400) \\ RMSLE &= 0.510 & (0.33) \end{aligned}$$

RMSLE Loss

- Specially useful for business cases where the underestimation of the target variable is not acceptable but overestimation can be tolerated.



Regression Analysis

SVM, KNN, Decision Trees

Recap: SVM Formulation


Minimize:

$$J(\mathbf{w}) = \frac{1}{2} \mathbf{w}^T \mathbf{w} + C \sum_{i=1}^N \xi_i$$



Subject to:

$$\forall_i, \quad y_i(\mathbf{w}^T \mathbf{x}_i + b) > 1 - \xi_i$$

- Solution is using QP solver
- **C** controls the relative importance of margin vs. training error



SVM Classification vs. Regression



For Regression:

- All samples should be within margin
- We want to minimize the margin
- Slack variables can be used to handle outliers


SVR Formulation

Minimize:

$$J(\mathbf{w}) = \frac{1}{2} \mathbf{w}^T \mathbf{w} + C \sum_{i=1}^N \xi_i$$

Subject to:

$$\forall_i, |y_i - \mathbf{w}^T \mathbf{x}_i| \leq \varepsilon + |\xi_i|$$




- FurtherReading:

Alex Smola and Bernhard Schölkopf, "A Tutorial on Support Vector Regression", Statistics and Computing 14: 199–222, 2004.

Non-linear SVR

- We can use the kernel trick for regression as well



sklearn.svm.SVR

```
from sklearn.svm import SVR
```

```
class SVR(*, kernel='rbf', degree=3, gamma='scale', coef0=0.0, tol=0.001,  
C=1.0, epsilon=0.1, shrinking=True, cache_size=200, verbose=False,  
max_iter=- 1)
```

- **kernel:** {'linear', 'poly', 'rbf', 'sigmoid'} or callable
- **degree:** int, default=3 // degree of polynomial kernel
- **gamma:** {'scale', 'auto'} or float // γ of poly, rbf or sigmoid kernels
- **C:** regularization parameter

For large datasets try using LinearSVR or SGDRegressor instead of SVR

<https://scikit-learn.org/stable/modules/generated/sklearn.svm.SVR.html>

k-NN Regression

- Key Idea: Similar inputs have similar predictions

Algorithm:

1. Find k-Nearest Neighbors of input \mathbf{x} : p_1, p_2, \dots, p_k .
2. Prediction for \mathbf{x} is a function of the y-values of these points:

$$\hat{y} = f(p_{1y}, p_{2y}, \dots, p_{ky})$$

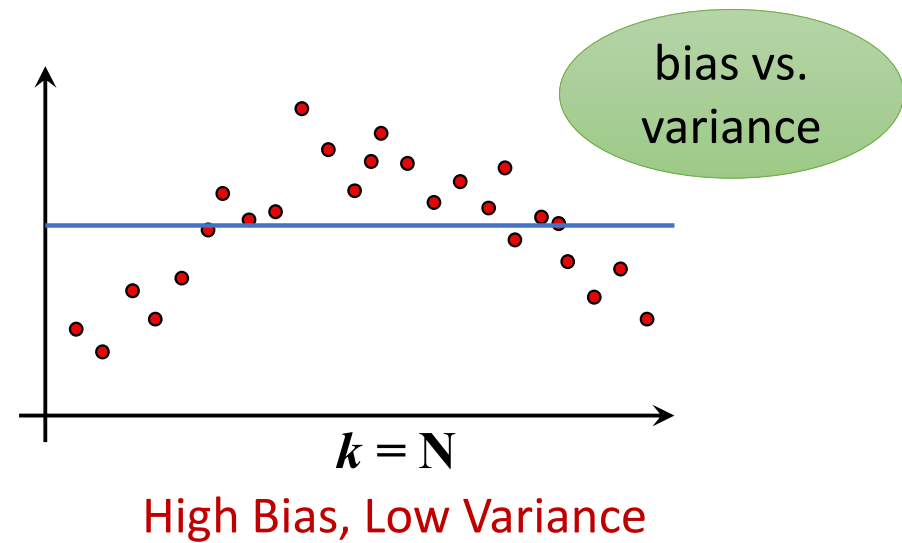


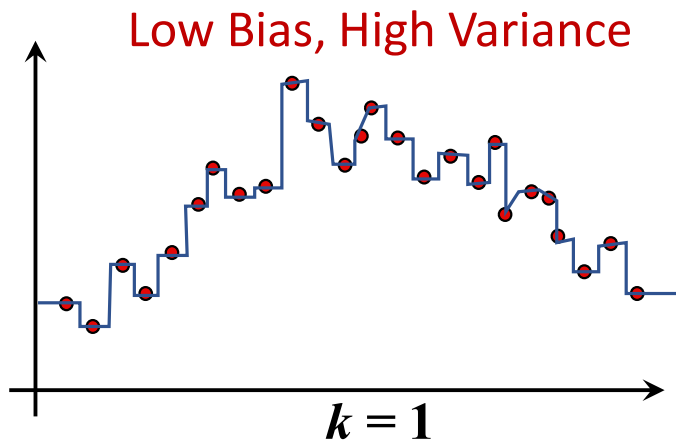
- Example 1: (simplest)

- $k = 1$; $\hat{y} = p_{1y}$

- Example 2:

- $k = 2$; $\hat{y} = (p_{1y} + p_{2y})/2$

Effect of k on prediction



kNN Regression: Comments

- Weighted Regression

- $\hat{y} = f(\mathbf{p}_{1y}, \mathbf{p}_{2y}, \dots, \mathbf{p}_{ky})$ is a weighted combination of $\mathbf{p}_{1y} \dots \mathbf{p}_{ky}$
- Uniform weights vs Weights based on distance
 - One could use different distance metrics
 - Weights can be inverse of distance or Gaussian

- Computing nearest Neighbors

- Exhaustive search
- KD-Tree, Ball-Tree

- Non-parametric Method

- k controls the bias-variance tradeoff

sklearn.neighbors.KNeighborsRegressor



```
class sklearn.neighbors.KNeighborsRegressor  
(n_neighbors=5, *, weights='uniform', algorithm='auto', leaf_size=30,  
p=2, metric='minkowski', metric_params=None, n_jobs=None)
```

- **n_neighbors:** Number of neighbors to use by default for k neighbors queries
- **weights:** {'uniform', 'distance'} or callable, default='uniform'
- **p:** int (for Minkowski), default=2
- **metric:** str or callable, default='minkowski'



<https://scikit-learn.org/stable/modules/generated/sklearn.neighbors.KNeighborsRegressor.html>

Recap: Decision Trees

1. Find the best feature (and threshold) to split the training data
 - Use an objective metric like Entropy to decide
2. Partition the training data as per the selected feature and threshold
3. For each partition, if the entropy is low, stop.
 - else, Repeat the first two steps for that partition



Example



- Compute sum of squared residuals for each split
- Choose the minimal one

Prediction and Split Criterion

- Value Predictor:
 - Common: Use the average value of training samples in a split
 - Minimum sum of squared residuals among all predictions
- Split Criterion
 - Find the variance of each potential split
 - Choose the split that minimizes total variance
 - Alternatively, maximize the reduction in variance

sklearn.tree.DecisionTreeRegressor

```
class sklearn.tree.DecisionTreeRegressor(*, criterion='squared_error',  
splitter='best', max_depth=None, min_samples_split=2, min_samples_leaf=1,  
min_weight_fraction_leaf=0.0, max_features=None, random_state=None,  
max_leaf_nodes=None, min_impurity_decrease=0.0, ccp_alpha=0.0)
```

- **criterion:** {"squared_error", "friedman_mse", "absolute_error", "poisson"}
- **max_depth:** int, default=None
- **min_samples_split:** int or float, default=2
- **min_samples_leaf:** int or float, default=1

<https://scikit-learn.org/stable/modules/generated/sklearn.tree.DecisionTreeRegressor.html>